

A DESIGN BIBLE FOR A WORLD-CLASS PRIVATE TRAVEL ATELIER

Travel Concierge.

*A luxury-for-less travel concierge: editorial in tone,
evidence-grounded in mind, cinematic in presence, and quietly
addictive in use.*

Volume: Vol. I · **Status:** Design strategy, pre-implementation · **Date:** May 2026

Audience: Founder · Design lead · Sonnet/Codex implementation pairs

Mantra: *Beauty without logic and logic without beauty are equally bad.*

This document does not modify code. It is the contract every future design pull request inherits. Pressure-test it, mark it up, and only then implement in the phased slices defined in Section 11.

CONTENTS

How to read this bible.

Each chapter stands alone. Read the executive summary first, then the chapter that matches the work you are about to do. The roadmap (§11) and guardrails (§12) bind every implementation prompt.

No.	Section
00	Executive summary & assumptions
01	Design diagnosis
02	Competitive synthesis
03	Brand and experience identity
04	Visual system
05	Motion and microinteraction system
06	AI Concierge flagship UX
07	Page-by-page product design plan
08	Card design system
09	Evidence and trust UX
10	Addictive personalization ideas
11	Implementation roadmap
12	Guardrails for future implementation prompts

13 First recommended design implementation slice

A Appendix · Glossary, success criteria, open questions

GUARDRAIL · This is a contract, not a wishlist

The 13 deliverables in this bible map 1:1 to the original brief. Section 11 (roadmap) sets which ones convert to PRs and in what order. Section 12 (guardrails) governs every PR prompt. Skipping either invalidates the bible.

SECTION 00

Executive summary & assumptions

A one-page contract before the bible begins.

What this document is

A complete design bible and phased redesign roadmap for the Travel Concierge product. It is opinionated, implementation-ready in spirit, and intentionally free of code. It is the artefact that future Sonnet and Codex prompts inherit.

What this document is not

It is not a generic visual mood board, a product spec, a backend plan, or a permission to start a sweeping redesign PR. The visual ambition lives inside a guardrail: never weaken intelligence, evidence, or speed for decoration.

Stated assumptions

- A1.** The current product already returns verified, addable Google Place cards in 3–8 seconds via the fast dynamic place search v1, and that pipeline is the source of truth for "addable" entities. Design must not regress that contract.
- A2.** A premium dark-mode foundation already shipped (PR #168) with deep navy + amber + cream tokens and Tailwind primitives. We extend, not replace.
- A3.** The codebase is Next.js + Tailwind on the frontend, with CSS variables already used for tokens. We can introduce additional tokens, but not a new CSS-in-JS system or a heavy animation library, without explicit budget approval.
- A4.** Backend, Supabase schema, and ranking logic are out of scope for this design bible. Design slices are frontend-only unless an explicit payload change is unavoidable, in which case it is split into its own backend PR first.
- A5.** The user is more sensitive to broken intelligence and broken speed than to visual roughness. Design is a force multiplier, not a justification for regressions.

Success criteria for the bible itself

- **Pressure-testable** — A reviewer can mark up any chapter and produce one concrete change, not a vague "I like it".
- **Phaseable** — The roadmap (§11) lets us ship the visual system in slices no larger than ~6 frontend files, each behind a budget gate.
- **Intelligence-safe** — Every page-level recommendation states what behaviour must not change, and how to detect a regression.
- **Evidence-honest** — No chapter encourages fabricating a vibe, view, award, or distance the backend has not verified.
- **Boutique, not bling** — No glow-on-hover, no color-soup gradients, no fake urgency, no childish gamification.

Headline recommendations

PRINCIPLE · The 6 things that change everything

1. **Adopt one tonal palette per surface.** Midnight ink for app shell, warm paper for trip artefacts, never both on one screen.
2. **Treat the Verified Place Card as the product's atomic unit.** Every other surface is a frame around it.
3. **Make evidence beautiful.** Trust marks, source counts, and caveats are first-class typographic citizens, not afterthought icons.
4. **Earn the AI Concierge's "intelligence" through restraint.** A composed thinking state and clean reason tags beat any animated avatar.
5. **Cinematic motion, deletable on demand.** 200–320ms easings, reduced-motion fallback baked in, no animation that delays first paint of cards.
6. **Phase 0 is tokens + card shell.** Nothing else. The first PR establishes the visual language; everything afterwards is a refactor against tokens.

How to use this bible with Sonnet/Codex

- **Always cite** — Every design PR prompt must reference the section(s) of this bible it is implementing, e.g. "Implements §11 Phase 1 + §8 Card Shell".
- **Never bundle** — A design PR may not touch backend, ranking, Supabase, or AI prompts. If it must, split.
- **Always token** — No raw hex on a component. If a token is missing, add the token first, then use it.
- **Always test the smoke path** — After every UI PR, manually run the spouse-friendly smoke test on mobile + desktop before merging.

SECTION 01

Design diagnosis

What the product risks becoming if visual work is handled piecemeal — and what we must protect on the way to a world-class concierge.

1.1 What the app risks becoming

The most realistic failure mode is not ugliness. It is the slow drift into a "competent travel SaaS dashboard" — clean, accessible, navigable, and emotionally inert. That product has no reason to exist next to Google Travel. The diagnosis below identifies the four drift vectors we must arrest before any decorative pass is allowed.

- **Drift 1 — “Card grid as wallpaper.”** — Every surface becomes a 3-column grid of equal-weight cards. The user cannot tell what is editorial, what is verified, what is AI-generated, and what is a saved idea. Information hierarchy collapses into a checkerboard.
- **Drift 2 — “Glow on hover.”** — The team mistakes neon hover states, gradient blobs, and soft shadows for luxury. The result feels like a 2021 crypto landing page, not a boutique concierge. Beauty-as-decoration without typographic restraint or content structure is the giveaway.
- **Drift 3 — “Chatbot UI inside a travel app.”** — The AI Concierge is rendered as a generic chat panel with avatar bubbles. The verified card — the actual product — is buried in a scrollable message thread. The conversational scaffold dominates the result.
- **Drift 4 — “Evidence theatre.”** — Trust signals are rendered as small green checkmarks tacked onto cards. They communicate “verified” the way a TLS padlock does — mute, commodity, defensive. They do not earn the price of the screen real estate.

1.2 What must be protected

PROTECT · Non-negotiables inherited from the product

Verified addable cards only. A card that can be added to a trip must come from a Google Place verification — never an editorial extract, never an LLM paraphrase, never a Tavily snippet. Visual design must communicate this rather than obscure it.

Evidence honesty. Reasons cite verifiable signals. If we lack evidence, we say so beautifully — we never invent a waterfront view, a Michelin mention, or a neighborhood claim.

Speed of return. Cards must paint in < 8s on first search, < 1s on pool/follow-up. No animation, image fetch, or font load may block first paint of card content.

Card identity stability. Follow-ups (“top 3”, “compare”, “best one”) reuse cards by identity key. Visual treatment must not re-issue, re-key, or re-fetch these cards.

The display ↔ supportingDetails ↔ whyPick contract. The three fields stay aligned. UI shall not invent a new field that fragments this contract.

1.3 Design debt to fix before polish

Polish on top of structural debt amplifies the debt. The following must be fixed inside the foundation phase before any decorative pass.

- **No global token contract.** — Some surfaces still ship light-era Tailwind utility classes. Tokens must be CSS variables, with a single source of truth in *globals.css*, and a naming convention that survives a re-skin.

- **Card shell variants are fragmented.** — Verified cards, AI Concierge cards, Saved Idea cards, and Itinerary cards each render through slightly different components. Reduce to one Card primitive with composable slots (media, identity, evidence, actions).
- **Trust marks are not first-class.** — Verified-by-Google, source count, evidence quality, and weak-evidence states are scattered across components. They need a single TrustStrip primitive.
- **Loading states are gray skeletons.** — Skeletons are correct but not concierge-grade. We need an editorial loading language: typeset placeholders, hairline pulses, deterministic order.
- **Empty states are dead.** — Empty states currently say “no results” instead of doing the concierge’s job — offering a thoughtful next step, a refinement chip, or an admission that the data is thin in this neighborhood.
- **Mobile is an afterthought.** — Several deep trip-detail pages are desktop-first. The luxury concierge feeling must work on a phone in a hotel lobby with one thumb.

1.4 Where design can amplify the AI Concierge instead of distract

Design earns its place when it makes the intelligence *legible*. The four amplification moves below cost nothing in latency and pay back every interaction.

- **Make the verified card the hero of the chat.** — The card, not the bubble, owns the visual weight. The conversation thread is a slim left rail of intent — the right side is the editorial result.
- **Render reasoning as quotation, not chrome.** — “Why this fits” is a typeset pull-quote with a small evidence tag, not an icon. It tells the user what the system noticed without exposing pipelines.
- **Communicate uncertainty as a feature.** — When evidence is thin, render a beautiful weak-evidence card with an explicit “verify when booking” line. Honesty becomes a luxury cue.
- **Compose the thinking state.** — Replace the spinner with a typeset “searching · verifying · composing” breadcrumb that matches the actual pipeline stages logged by *fast_dynamic_place_search.timing*.

SECTION 02

Competitive synthesis

What to learn, what to reject, and what we can do better than every reference product.

Every reference below is studied for one thing: the smallest pattern we can borrow without inheriting its limitations. Generic praise (“Airbnb is clean”) is forbidden in this chapter.

Airbnb

What to learn

- Information density per card is exceptional. Photo-first hierarchy with price + rating + neighborhood as a tight metadata strip.
- Map ↔ list co-presence on desktop with a sticky-pan filter that does not requery on tiny pan deltas.
- Calendar UI handles flex dates gracefully without modal hell.

What to reject

- The “Rare Finds / Luxe” category split is hand-curated marketing. We will not copy categories that imply editorial human review we do not have.
- The hero gallery on a listing page is so dominant it crowds out trust evidence. We will keep media restrained on verified cards.
- Recent dark-mode and Stays/Experiences toggles feel like SaaS chrome on top of editorial photography. We avoid that collision by committing to one tonal mode per surface.

How Travel Concierge wins

Borrow the metadata strip rhythm. Reject the photo-as-everything hierarchy. Beat them on evidence: where Airbnb shows star ratings, we show structured “why this fits” with sources.

Expedia

What to learn

- Multi-leg flight + hotel + car bundling logic is the strongest in the industry. The data model exposes price-by-component cleanly.
- Filter chips persist across a session and restore from URL state.

What to reject

- The visual language is loud and promotional — banners, sale badges, urgency timers. We will not import any of it.
- Comparison views become spreadsheet-grade dense. Boutique users bounce.

How Travel Concierge wins

Borrow the URL-stateful filter pattern for Explore. Reject the ad-banner aesthetic outright. Beat them on calm: the absence of “only 3 left!” is the design statement.

Booking.com

What to learn

- Price transparency is excellent — total-with-fees is shown by default in many regions.
- Map clusters degrade gracefully at low zoom.

What to reject

- Dark patterns: scarcity messaging, “85 people viewing now”, color-coded urgency. Forbidden.
- Hyper-saturated blues and yellows that telegraph “discount aggregator”.

How Travel Concierge wins

Borrow the all-in price honesty norm and the map cluster behavior. Beat them on dignity: no scarcity theatre, ever.

Tripadvisor

What to learn

- User reviews aggregate at scale and make “consensus quality” a useful signal even when individual reviews are noisy.
- “Things to do” taxonomy is a useful POI ontology.

What to reject

- Visual identity is undifferentiated — green-on-white, ad-driven.
- Editorial voice has been worn down to advertorial. Trust is leaking.

How Travel Concierge wins

Borrow the consensus aggregation idea, but render it as our own evidence ladder (see §9). Reject the advertorial voice. Beat them on voice: a concierge that sounds like a person, not a feed.

Google Travel & Google AI travel planning

What to learn

- Map-first navigation; instant zoom-to-pin; the place panel is the highest-fidelity POI surface in production.
- AI itinerary surfaces (Bard / Gemini) cite sources inline and let users edit a structured plan.

What to reject

- Surfaces are emotionally cold. The product is a utility, not an experience.
- Cards have no “personality” — they are extractive panels of GMB data.
- AI-generated plans collapse into bulleted text without addable identity.

How Travel Concierge wins

Borrow the place-panel fidelity for our verified card detail. Reject the cold utility tone. Beat them on warmth: the same data, dressed for evening.

Kayak / Hopper / Skyscanner

What to learn

- Calendar heatmaps, price-prediction indicators, and flexible-month grids.
- Multi-airport selection patterns and inline error messages.

What to reject

- Loud color-coded prediction badges that turn the screen into a casino.
- In Hopper specifically, the mascot-driven UI is too whimsical for our brand.

How Travel Concierge wins

Borrow the calendar-heatmap mental model for trip-date suggestions. Reject the casino color palette. Beat them on calm: a single typographic price-trend line.

Mr & Mrs Smith / Belmond / boutique luxury sites

What to learn

- Editorial photography commitment: full-bleed images shot for the brand, never stock. Type pairing of a serif display with a clean sans body.
- “Insider” voice — first-person, knowing, uncluttered.

What to reject

- Booking flows are often anaemic on these sites. They optimize for awe, not for trip planning.
- Mobile experiences are sometimes sacrificed to art-directed desktop hero sequences.

How Travel Concierge wins

Borrow the type pairing, the insider voice, the photographic commitment. Beat them on utility: we will not let the editorial layer slow trip planning by a single second.

Awwwards luxury hospitality & AI-first planners

What to learn

- Custom cursor work, scroll-driven storytelling, restrained color systems, serif-display typography.
- AI-first planners (e.g., Layla, Wonderplan, Mindtrip) experiment with inline card-in-chat patterns.

What to reject

- Most are demo-grade. Performance, accessibility, and content depth are frequently sacrificed for the showreel.
- AI planners often invent reasons. They are exactly what we must not become.

How Travel Concierge wins

Borrow the typographic discipline and the inline card-in-chat shape. Reject every performance-sacrificing flourish. Beat them on substance: we will be the AI planner that *cites its sources*.

2.9 Patterns we synthesize across the field

- **Map ↔ list parity.** — Adopted from Airbnb, Google. We treat the map as a first-class twin of the list, not a popup. Filter changes update both atomically.
- **Inline citations.** — Adopted from Google AI; we show source provenance under every reason chip.
- **Editorial typography.** — Adopted from Mr & Mrs Smith and Awwwards; serif display + restrained sans body, never both serif.
- **All-in pricing posture.** — Adopted from Booking; price honesty becomes a brand signal.
- **Calendar heatmap mental model.** — Adopted from Kayak; we use it sparingly to suggest cheaper trip windows when we have provider data, never as urgency.

SECTION 03

Brand and experience identity

The product personality, voice, and emotional bar — in concrete, testable terms.

3.1 Emotional tone

Walking into a small, candle-lit travel atelier on a side street. The owner remembers your last trip, has the right book on the shelf, and never raises their voice. Confidence without showmanship. Warmth without sweetness. Speed without urgency.

3.2 Product personality

Trait	Means	Does not mean
Boutique	Hand-feel; restraint; specific choices	Niche, exclusionary, esoteric
Editorial	Voice; typographic care; long-form when warranted	Magazine pastiche, cover lines, splash images on every screen
Concierge	Anticipatory; remembers context; offers next-best-action	Servile, fawning, scripted
Honest	Cites evidence; admits weak signal; shows uncertainty	Self-flagellating; over-apologizing; “as an AI I cannot...”
Quiet	Calm motion; no urgency; no shouting	Sleepy, slow, lacking opinion
Generous	“Worth-the-splurge” + “luxury-for-less” equally weighted	Couponing, deal-bait, fake markdowns

3.3 Visual language headline

- **Tonal duality.** — Two surfaces only: **Midnight Ink** for the app shell, concierge thinking, and the AI surface; **Warm Paper** for itineraries, saved scrapbooks, and shareable trip artefacts. Never both at once.
- **Type as identity.** — A display serif with personality (e.g., Tiempos Headline, Editorial New, GT Super Display, or self-hosted equivalents) paired with a humanist sans (e.g., Söhne, Inter, or system-ui) for body. No third family.
- **Photograph or no image.** — Where we can show a real photograph, we go full-bleed and edge-to-content. Where we cannot, we never put a stock image on screen — we typeset.
- **Hairline geometry.** — Borders, dividers, table rules sit at 0.4–0.6px. The grid is felt, not seen.
- **Light as lacquer, not as glow.** — A single subtle radial highlight on raised surfaces in dark mode. No glow on hover; the cursor is the highlight.

3.4 Interaction language

- **Direct manipulation over modal forms.** — Drag a card to a day. Tap a chip to refine. Long-press to compare. Modals only for legally distinct actions (sign out, delete trip).
- **Persistence is a vibe.** — Filters, scroll positions, and prompt history survive page transitions. The concierge does not forget what you asked five seconds ago.

- **Two-finger generosity on mobile.** — Every primary action is reachable by the right thumb. Secondary actions live in a tap-to-reveal sheet, not in a hamburger.
- **Keyboard for power users.** — `k` = open command bar; `■K` opens AI Concierge from anywhere; `/` focuses search; `esc` closes any drawer. Keyboard shortcuts are *complementary*, not required.

3.5 Content voice

A first-person concierge who has been to the city, knows what the user means when they say “tapas bar,” and is too dignified to upsell. We write as if a real person is composing each line.

Use	Avoid
"A stronger tapas/small-plates match than the cocktail bar around here."	"Highly rated!" "Top pick! Don't miss out!"
"Worth booking ahead on weekends — it fills early."	"Limited availability — only 3 spots left!"
"We could not verify the waterfront view; ask when booking."	"Stunning waterfront views!" (when unverified)
"You saved three hidden gems in River North."	"You've unlocked the River North expert badge ■"
"This is closer to your hotel than the others by ~12 minutes on foot."	"OMG so close!! ■"

3.6 Trust language

- **“Verified by Google.”** — Used only when we have a stable Google Place ID and OPERATIONAL status. Never as decoration.
- **“Reviews aggregated from N sources.”** — A small numeric strip with a link to expand the source list. Never a fake aggregate score.
- **“Confidence: high / medium / low” mapped to evidence count.** — High = 3+ corroborating sources; Medium = 2; Low = 1 + verified place. The language is plain English, not a numeric percentile.
- **“We could not verify...”** — A first-class sentence pattern. It appears as a typeset caveat under any reason that includes a constraint we did not confirm.

3.7 Luxury-for-less positioning

The brand is not “cheap luxury”. It is *known luxury*: we know which splurges return value and which do not. The product surfaces this through two recurring chips on cards.

- **“Worth the splurge.”** — Applied only when the verified place has a meaningfully higher price band and corroborating evidence of a notable experience. Never on a generic mid-tier place.
- **“Luxury for less.”** — Applied when a verified place delivers a high-end experience at a notably lower price band than peers in the same neighborhood. Both chips use the same sandstone gold accent — siblings, not rivals.

3.8 What the app must never feel like

REJECT · Anti-patterns the brand forbids

OTA aggregator. No "Best price guarantee", no countdown timers, no "Lowest prices in 24 hours" badges.

SaaS dashboard. No KPI tiles, sparklines, "completion percentage" rings, or productivity-app empty states.

Crypto landing page. No animated gradient meshes, no neon hover glows, no glitch type, no "ETA: live data" chips.

Generic chatbot. No animated thinking dots inside speech bubbles, no avatar with cute facial expressions, no "Hi! I'm your AI assistant!" opener.

Gamified app. No streaks, no XP, no badges, no level-up modals, no confetti, ever.

Stock-photo magazine. No generic "sunset over a city" hero. Either we have a real verified-place photo or we typeset.

SECTION 04

Visual system

A complete token-level visual direction: dark-primary, paper-secondary, with motion-aware surfaces and accessible contrast guarantees.

4.1 Dark mode (primary system)

Dark mode is the canonical canvas. It hosts the AI Concierge, Explore, the shell, and any surface where the product is *composing* for the user. It is built from a five-step ink ladder, not a single black.


Midnight Ink #0B1320 <i>page background</i>	Onyx Velvet #0F1A2C <i>surface</i>	Carbon Mist #1A2538 <i>raised surface</i>
Pen Stroke #22324A <i>divider</i>	Sandstone Gold #E0B888 <i>primary accent</i>	Ember Brass #C5944D <i>deep accent</i>
Pearl Cream #F2EBDD <i>headline text</i>	Cream #E8E2D4 <i>body text</i>	Mist #9AA4B2 <i>subdued text</i>
Verified Sage #88A899 <i>trust signal</i>	Caution Amber #E8B26B <i>soft caution</i>	Whisper Coral #D88478 <i>gentle warning</i>

Every surface uses one ink for background, one ink one step lighter for the card, and one ink one step lighter again for the raised state. No more, no less. Cards on cards on cards is the killer of dark UI.

4.2 Light mode (secondary system)

Light mode is the *artefact* mode — trip itineraries that someone might print or screenshot to share, the saved-ideas scrapbook, and the user profile. It is warm paper, not stark white.

Warm Paper #FAF7F0 <i>page background</i>	Bone #F1ECE0 <i>surface</i>	Linen #E6DECB <i>raised surface</i>
--	--	--



4.3 Color roles, not random colors

Role token	Dark mode	Light mode	Used for
--surface-0	Midnight Ink	Warm Paper	Page canvas
--surface-1	Onyx Velvet	Bone	Default cards
--surface-2	Carbon Mist	Linen	Raised cards / drawers
--divider	Pen Stroke	Hairline	Hairline borders, table rules
--text-strong	Pearl Cream	Ink Paper	Headlines
--text-body	Cream	Slate	Body copy
--text-mute	Mist	Slate@70%	Captions, timestamps
--accent	Sandstone Gold	Brass	Primary CTAs, "Add to trip"
--accent-deep	Ember Brass	Brass-deep	Hover/pressed accent
--trust	Verified Sage	Sage	"Verified by Google" mark
--caution	Caution Amber	Amber	Weak-evidence chip
--warn	Whisper Coral	Coral	Booking caveat, blocked action

4.4 Typography roles

Two families. One serif display, one humanist sans. A monospace for evidence/code/IDs only. No third family. Sizes climb a 1.25 modular scale anchored at 16px body.

Role	Family	Size / Leading	Use
Display XL	Serif Display	64 / 68	Cover hero, landing
Display L	Serif Display	44 / 50	Page hero / section title
Display M	Serif Display	32 / 38	Card detail title

Display S	Serif Display	24 / 30	Section title in card
Body L	Sans	18 / 28	Lead paragraph
Body	Sans	15 / 24	Default body
Body S	Sans	13 / 20	Card metadata, list rows
Caption	Sans	12 / 16	Trust marks, timestamps
Overline	Sans (caps, +tracking)	10 / 14	Section labels
Mono	Mono	12 / 16	Identity keys, debug only
Quote	Serif Italic	18 / 28	Concierge reasoning quotes

4.5 Spacing rhythm

A single 4px base unit, with semantic tokens above it. Hairlines and borders sit on a 1px grid; everything else snaps to 4.

Token	Pixels	Usage
--space-1	4	Sub-icon padding
--space-2	8	Default chip padding, inline gap
--space-3	12	Card metadata gap
--space-4	16	Card padding (compact)
--space-5	20	Card padding (default)
--space-6	24	Section gap
--space-8	32	Heading-to-content gap
--space-10	40	Card-to-card gap on Explore
--space-12	48	Hero block padding
--space-16	64	Page-edge gutter (desktop)

4.6 Card elevation system

In dark mode, elevation is achieved by stepping up the ink ladder, never by increasing shadow softness. In light mode, elevation is achieved with hairline borders and a 2%-luminance surface lift, not drop shadows.

Level	Dark mode	Light mode	Use
e0	Surface 0	Paper	Page canvas
e1	Surface 1 + 0.4 px border	Bone + 0.4 px border	Default card
e2	Surface 2 + 1 px border	Linen + 1 px border	Hovered/active card

e3	Surface 2 + soft 8&thinspx shadow @ 18% black + 8&thinspx shadow @ 6% Drawer, modal
e4	Surface 2 + 16&thinspx shadow @ 24% black + 16&thinspx shadow @ 10% Floating composer

4.7 Borders, shadows, glass, texture, gradients

- **Borders.** — Always 0.4–1px hairlines. Token: `--divider`. Border-radius scale: 4 (chip), 8 (card), 12 (drawer), 16 (modal). No 24+; we are not a soft-toy app.
- **Shadows.** — Used sparingly, only on e3/e4. Two-stop: a tight 1px shadow for crispness, plus a soft 12–24 px shadow for atmosphere. Shadows are warmer (slightly brass-tinted) in dark mode.
- **Glass.** — One use only: the AI Concierge composer floating over a destination photo on landing. *backdrop-filter: blur(22px) saturate(140%)*. Forbidden everywhere else — it kills mobile performance and hides text.
- **Texture.** — A 1%-opacity film grain layer is permitted on the landing and login page only, baked into a single SVG, no JS animation.
- **Gradients.** — Two permitted: a vertical ink-fade behind the AI Concierge thinking state, and a 2-stop sandstone-to-brass on the primary CTA hover. Banned: 4+ stop conic gradients, neon edges, animated mesh.

4.8 Iconography

Single-weight 1.5px stroke, 24px grid, rounded joins, no fills except for filled "saved" states. Source: Lucide or Phosphor; never mix two icon libraries on one screen. Icons never carry semantic meaning alone — always paired with a label or aria-label.

4.9 Map styling direction

- **Two map themes, both ours.** — Dark map: deep ink with brass roads and sage parks. Paper map: cream with umber roads and sage parks. Both lower-saturation than default Google.
- **Pins as identity, not as decoration.** — A verified place pin is a small filled circle in sandstone gold; a hovered pin gains a 1.5x ring. Saved places get a hairline ring around the pin.
- **Cluster typography.** — Cluster counts are typeset in the same humanist sans as the app, never a system default.
- **Density discipline.** — Show at most 30 pins on screen. Beyond that, cluster. Never paint a swarm.

4.10 Image treatment

- **Aspect ratios.** — 4:5 portrait for verified place hero; 3:2 landscape for Explore covers; 1:1 square for Saved Idea thumbnails. No 16:9 unless we are showing a video clip we control.

- **Edges.** — All images sit inside the card; we never bleed them to the page edge except on the cover/landing hero.
- **Treatment.** — A subtle 4% midnight-ink overlay in dark mode for legibility of overlaid metadata; never a vignette, never a photo filter that distorts color.
- **Fallbacks.** — No stock photos, ever. If we have no photo, we typeset. The card surfaces a monogram letter in serif display + a tonal background derived from the place type.
- **Loading.** — Images load progressively: tonal background → blurred low-res → full. Card content paints before the image — the user reads the reason first.

4.11 Empty, loading, and error states

- **Empty: a sentence, then a refinement.** — “Nothing matches %s in this neighborhood yet. Try widening to %s, or ask for a different cuisine.” Always include two clickable refinements derived from the prior search.
- **Loading: a typeset breadcrumb.** — *Searching · Verifying · Composing*, with the active stage in pearl cream and inactive stages in mist. Progress is real (driven by *fast_dynamic_place_search.timing*), not theatre.
- **Error: name the cause, offer a path.** — “The Google Places verifier is unavailable. We will retry; meanwhile, here are your saved places.” No stack traces, no “Oops!” cartoons.
- **Skeletons: typed, not gray.** — Card skeletons render the actual card geometry with hairline outlines; the place name is a 60%-width typeset bar, not a flat rectangle.

4.12 Mobile responsiveness principles

- **Three breakpoints only.** — Phone (< 640px), tablet (640–1024), desktop (≥ 1024). No between-breakpoint adjustments.
- **Right-thumb reachability.** — Primary actions sit in the bottom 25% on phone, never in a fixed top app-bar pill.
- **AI Concierge as a sheet, not a page, on mobile.** — It opens as a 90vh bottom sheet that can be dragged closed. The composer is always visible.
- **Map → list precedence.** — On phone, list is default; map is reachable through a single chip. We do not split the screen 50/50 on phones.

4.13 Accessibility constraints

- **Contrast.** — All text passes WCAG 2.2 AA at the rendered size. Pearl Cream on Midnight Ink ≈ 13.4:1; Cream on Onyx ≈ 11.2:1. We test every accent on both surfaces.
- **Focus.** — A 2px sandstone-gold focus ring with 2px offset on every interactive element. Never *outline: none*.
- **Motion.** — *prefers-reduced-motion* disables all entrance animations and reduces durations to ≤ 80ms.

- **Hit area.** — 44×44px minimum on touch. Chips have invisible padding to reach this without inflating visual size.
- **Text resizing.** — Layout survives 200% browser zoom without horizontal scroll.
- **Color is never the only signal.** — Trust marks pair sage with a label; weak-evidence pairs amber with the phrase “We could not verify...”.

SECTION 05

Motion and microinteraction system

Cinematic motion that respects performance, reduced-motion preferences, and the speed of the intelligence pipeline.

5.1 Motion principles

- **Motion explains, never decorates.** — Every motion answers a question: where did this come from, where did it go, why is it different now? Motion that does not answer one of these is removed.
- **Easing is a brand asset.** — A single ease-out-quad on entrance and ease-in-quad on exit. No bouncy springs, no overshoot, no elastic. Boutique motion is contained.
- **Durations are short and rhythmic.** — Micro: 120ms (chip press, focus). Default: 200ms (card hover, drawer open). Spatial: 320ms (page transition, sheet enter). Nothing longer than 400ms.
- **Motion never blocks first paint.** — Card content renders immediately at final position, then fades from 0.6 → 1 opacity. We never animate *from offscreen* → *onscreen* for primary content.
- **Reduced-motion is a real product.** — *prefers-reduced-motion* swaps every animation for a 60ms opacity fade. Every component must be tested in this mode.

5.2 Page transitions

Transition	From → To	Treatment	Duration
Trip enter	Trips list → Trip detail	Trip cover image cross-fades; sidebar persists	320 ms
Concierge open	Any page → AI Concierge drawer	Drawer slides in from right (desktop) / up (mobile); page content fades out	200ms to 60% 200ms to 60% 200ms to 60% 200ms to 60%
Tab switch	Explore tabs	Content cross-fades 200 ms; tab indicator slides 220 ms	200 ms
Auth → app	Login → Dashboard	Cinematic dim of background image; serif-heading entrance	600ms
Modal open	Any → modal	Backdrop fades in 200 ms; modal scales 0.96 → 1 with 240ms ease-out	240ms

5.3 Card entrance

When a verified card list returns from the AI Concierge, cards fade in with a 40ms stagger between siblings, in score order. The stagger is capped: after the 6th card we stop staggering, so the 12th result is not delayed by 480 ms. Cards that come from a cache or pool hit do **not** stagger — they paint at once, signalling speed.

PRINCIPLE · What happens when the pipeline is faster than the animation

On a pool/follow-up hit, total backend time can be < 1s. We must not invent a 600ms thinking animation to make the system look "smart". When data arrives in < 200ms, render immediately with no staggered entrance. Speed is the luxury cue.

5.4 AI Concierge thinking, searching, verifying states

- **Composed breadcrumb.** — *Searching · Verifying · Composing*. Active stage glows gently (opacity 1.0 + 1 px pearl underline); inactive stages at opacity 0.4. Stage transitions are driven by real timing events from *fast_dynamic_place_search.timing*.

- **Cursor pulse, not spinner.** — A single 1px pearl-cream caret blinks at the end of the active stage label. No spinner, no loading bar, no animated dots.
- **Verification micro-moment.** — When a Google Place verifies, a small sage hairline draws across the bottom of the appearing card in 220ms. This is the only "verification animation" we ship.
- **Hard cap.** — If the pipeline exceeds 12s, the breadcrumb adds a fourth stage, "Stretching the search", and surfaces a "show what you have so far" chip. We never let the user wait silently.

5.5 Save / Add-to-trip moments

- **Add-to-trip.** — The card briefly slides 6px right while a sandstone hairline pulses across it. A toast does NOT appear; instead, the trip's tab indicator in the sidebar gains a 1.5x dot for 1.6s.
- **Save (heart) toggle.** — The icon transitions stroke→fill in 160ms with a 4% scale pulse. No confetti, no toast.
- **Move to a day.** — A drag handle is always visible on itinerary cards on desktop. A long-press starts drag on mobile. Drop zones gain a sandstone hairline on hover; the whole row is the drop target, not a tiny gap between cards.

5.6 Timeline transitions

- **Day reorder.** — Re-ordering items animates the surrounding rows 200ms with FLIP. No vertical scroll jump.
- **Travel-time hint reveal.** — When a hint inserts (e.g., "12 min walk"), it expands from height 0 → computed height in 220ms. The hint is hairline + caption type, never a full card.
- **Day collapse.** — Days collapse to a single row that shows the day's "headline pick" + count. Expansion uses the same 220ms transition.

5.7 Map ↔ card interactions

- **Hover-to-pin.** — Hovering a card highlights its pin (1.5x ring) and dims unrelated pins to 0.4 opacity. Hovering a pin highlights its card by 1px sandstone border — the list does NOT auto-scroll.
- **Click-to-fly.** — Clicking a pin flies the map to the pin in 320ms with ease-out-cubic and opens the card's detail drawer. Clicking the same pin twice closes the drawer.
- **Pan locks the requery.** — A small "search this area" chip appears in the top-center of the map after a 300px pan. We never auto-requery on pan.

5.8 Hover and tap states

- **Cards.** — Hover lifts elevation e1 → e2 (border becomes 1px sandstone @ 30%). No scale, no shadow puff, no glow.

- **Buttons.** — Primary button hover: background goes from sandstone to ember-brass over 120ms. Secondary: hairline border thickens 0.5 → 1px. Pressed: a 1px inner shadow, no scale.
- **Chips.** — Hover: 1px sandstone outline. Active: filled sandstone with ink text. No "wiggle" or "bounce".

5.9 Forbidden animations

REJECT · Animations that harm the brand or the pipeline

Parallax scrolling. Slows everything; not boutique.

Glow pulses on hover. Generic SaaS aesthetic.

Animated gradient meshes. Crypto-landing aesthetic.

Bouncing/elasticity springs. Not boutique.

Confetti, fireworks, “success” bursts. Childish.

Auto-playing video backgrounds. Slow on mobile, dishonest evidence.

Scroll-jacking. Removes the user’s control of pacing.

Letter-by-letter typewriter effects on AI replies. Adds latency, feels like a chatbot, undermines our card-first hierarchy.

5.10 Reduced-motion fallback

@media (prefers-reduced-motion: reduce) applies the following overrides globally: all transform animations disabled; all opacity transitions capped at 60ms; all stagger removed; map fly-to becomes a hard cut. The product still feels intentional; it just stops moving.

SECTION 06

AI Concierge flagship UX

The flagship surface. Treats verified cards as the hero and the conversation as a slim editorial rail.

The AI Concierge must feel like the most intelligent surface in the product without exposing pipeline internals, debug payloads, or cute chatbot tropes. The design move is structural: invert the dominance from chat-bubble-with-cards to cards-with-conversation-rail.

6.1 Layout

Two-column layout on desktop ($\geq 1024\text{px}$); single-column sheet on mobile.

Region	Width	Contents
Conversation rail (left)	320 px	Prompt history as small typeset turns, intent chip, "show evidence" toggle, memory pill ("p
Result canvas (right)	fills	Verified card grid (3 columns desktop, 2 tablet, 1 mobile), with the active card promotable
Composer (sticky)	full	Always visible on the bottom 88 px of the result canvas. Houses the prompt input,

6.2 Prompt composer

- **A typeset textarea, not a chat input.** — Sans 16px, 4-line max, autoresizes. Placeholder rotates between three concrete examples per session: "sushi with a waterfront view", "tapas bar that is not too loud", "the ones I'd actually take my mother to". Placeholders never advertise capabilities.
- **Composer chips.** — Three contextual chips above the input: **Refine**, **Compare**, **More options**. They post the prompt with the right intent so the user does not type the magic phrase.
- **Submit on enter, newline on shift-enter.** — Standard. The send affordance is a 24px outlined arrow in sandstone, not a "Send" button.
- **Token budget hint.** — A faint counter appears only when the prompt exceeds 240 characters: "Long prompt — we will summarize." Honest, restrained.

6.3 Verified card result layout

PRINCIPLE · Card-first hierarchy

The conversation rail summarizes the user's ask in one line: *"Tapas bars in West Loop, evening, dinner-date feel."* Below, the result canvas opens directly into the card grid — **no AI reply paragraph**. The prose belongs in the card's "why this fits" section, not above the grid. This is the single most opinionated decision in this bible.

- **No "Here are some options for you" preamble.** — Preambles are a chatbot tell. They cost an extra line of vertical space, push the verified cards below the fold, and add nothing the cards do not say.
- **Result count is editorial.** — "Six places that fit." with the count rendered as a small Display S serif. Not "I found 6 results."
- **A "Why these six" link.** — Below the count, a small caption: *How we picked these — tap to see*. Expands a 1-paragraph plain-language summary citing the constraints we detected (cuisine, vibe, neighborhood).

Honest, brief.

6.4 Evidence and trust indicators on cards

See §9 for the full evidence chapter. Inside the AI Concierge result canvas, the trust strip on each card is identical to the trust strip on the rest of the product. The Concierge is not allowed to invent its own trust language.

6.5 Follow-up chips

Below the result grid, a row of follow-up chips. They are not generic; they are derived from the parsed intent (cuisine, vibe, constraint, neighborhood) and from the cards just returned.

- **Refinement chips.** — "Quieter", "More upscale", "Walking distance from my hotel", "Better wine list". Each one mutates the prompt; we never invent a constraint we did not detect.
- **Compare chip.** — "Compare top 3" routes to the existing refine_previous flow without re-issuing card identity keys.
- **Negate chip.** — "Not these" lets the user dismiss the current set and re-search with a negative-constraint hint.

6.6 Conversational memory indicators

- **Memory pill.** — Top-left of the rail: a tiny pill summarizing the active context — "*Chicago · May 2026 · anniversary trip*". Tapping the pill opens an edit drawer. The pill is the only place memory is exposed.
- **Edit, do not deny.** — If the user wants to change the destination, we let them edit the pill. We never "auto-clear" memory or pretend we forgot.
- **Forget toggle.** — A "start fresh" link in the pill drawer wipes context cleanly. We do not hide this.

6.7 Weak-evidence handling

When evidence is thin (single source, place verified but key constraint not corroborated), the card surfaces a *weak-evidence chip* in caution amber, and the "why this fits" line begins with "*Likely fits — we could not verify...*". We never hide the weak-evidence card; we present it honestly.

OPPORTUNITY · Weak evidence is a feature, not a failure

Other AI planners hide their uncertainty by inventing prose. Travel Concierge turns uncertainty into a luxury cue: a beautifully typeset caveat says we looked, we did not find corroboration, and we will help the user verify on booking. Honesty becomes the brand.

6.8 No-results handling

When zero verified places match, we never show a “No results found” state. We show a typeset, three-paragraph editorial empty state:

"Nothing in West Loop matched 'tapas bar' tonight that we are confident about. The closest fits we can find are in River North, ~12min away. You can broaden the area, change the cuisine, or open the search to tomorrow night."

Three refinement chips appear under the empty state, derived from the parsed intent.

6.9 Card add flow

- **Single-tap add.** — A primary "Add to trip" button on every verified card. If the user has more than one trip, a small popover lets them choose — never a full modal.
- **Default trip.** — The most-recently-touched trip is the default. The popover offers other trips as a list, with a "Save without trip" option.
- **Acknowledgement.** — See §5.5. The card briefly slides right with a sandstone pulse; no toast.

6.10 Compare / refine flow

- **Top-3 compare.** — "Compare top 3" promotes three cards to a horizontal scroll lane with the evidence rows aligned: cuisine match, vibe match, source count, walking distance from hotel. Each row is a typeset list, not a table chrome.
- **Refine without re-fetching.** — A refinement chip ("quieter", "better wine list") that we can satisfy from existing card metadata never triggers a new search; we re-rank the existing pool. The composer shows a hairline pulse rather than the full thinking breadcrumb — the user feels the speed.

6.11 Mobile behavior

- **Concierge as a 90vh sheet.** — Opened from any screen via the floating composer pill or ■K equivalent. Drag-down dismiss; the composer is always visible.
- **Conversation rail collapses to a header.** — Memory pill + intent line live in the sheet header; full history opens via a small "history" link — not a hamburger.
- **Single-column card stack.** — Cards stack at 100% width, with horizontal swipe to compare neighbors. The add-to-trip button stays sticky inside the sheet.

6.12 Making it feel intelligent without exposing internals

GUARDRAIL · Intelligence cues we ship vs. ones we forbid

Ship: the parsed-intent line ("Tapas bars in West Loop, evening, dinner-date feel."), the composed thinking breadcrumb, the dynamic "why this fits" line, the weak-evidence chip, the source count, the refine chips derived from constraints we actually parsed.

Forbid: animated thinking dots, simulated typing, a chatbot avatar, a "thought process" panel showing chain-of-thought, exposing model name or token counts, exposing the raw search query we sent to Google, exposing place IDs or identity keys.

SECTION 07

Page-by-page product design plan

Every major surface, with intent, hierarchy, components, states, risks, and what must not change yet.

For each surface we list: *user intent, visual goal, information hierarchy, key components, interactions, empty/loading/error, risks, implementation notes, what not to change yet.*

7.1 Unauthenticated landing

User intent

Decide in 8 seconds whether this product is for me.

Visual goal

A cinematic hero, not a SaaS landing. Single full-bleed destination photograph, serif display headline, single CTA.

Information hierarchy

- 1. Headline statement of brand promise (serif Display XL).
- 2. One-line value proposition (sans body).
- 3. Single CTA: "Plan your trip." (no secondary CTA above the fold).
- 4. Three-card editorial proof strip below the fold (no logos wall).
- 5. AI Concierge inline demo — a real, throttled demo prompt.

Key components

- Cinematic hero with photo + 4% ink overlay + 1% grain.
- Composer pill (glass, the only glass surface in the product).
- Editorial proof strip: three real verified-place cards.
- Footer with restraint — no eight-column link grid.

Interactions

- Composer pill in the hero accepts a real prompt; submitting routes to sign-up with the prompt prefilled in the AI Concierge.
- Scroll triggers a 600ms ink fade behind the proof strip.

Empty / loading / error

No empty state. Loading: only the proof strip awaits an image; the rest of the page paints immediately.

Risks

- Glass + photo can hurt mobile FCP. Mitigate by serving an LQIP and capping blur radius on phone.
- A real demo composer can be misused. Throttle by IP and by prompt length.

Implementation notes

- Photo asset must be licensed; no Unsplash placeholders shipped to prod.
- Demo composer goes to the existing AI route with a "demo" flag; no new API.

GUARDRAIL · What not to change yet

Do not redesign the auth flow yet; the demo only routes to existing sign-up.

Do not introduce a video hero before mobile FCP budgets are set.

7.2 Login / signup

User intent

Get in fast, feel like I have arrived somewhere nice.

Visual goal

A still hero (not a moving one), a single column form, identical type and tokens to the rest of the product.

Information hierarchy

- 1. Brand wordmark, 24px, top-left.
- 2. Serif Display L: "Welcome back."
- 3. Email + password fields, OAuth buttons stacked below.
- 4. Switch to sign-up link, plain underline, no card.

Key components

- Cinematic still photo (swappable for one of three curated destinations).
- Single-column form, 360px width, 24px field gap.
- Inline error: hairline coral border + sentence below the field.

Interactions

- Submit on enter; OAuth buttons go to provider modal directly.
- No "remember me" checkbox — we use long-lived sessions.

Empty / loading / error

Loading: button label changes to "Signing you in"; no spinner. Error: hairline coral, plain English message ("That password is wrong" not "Authentication failed").

Risks

- Password reset flow lives outside this slice; do not bundle.

Implementation notes

- Reuse existing auth API; no payload changes.

GUARDRAIL · What not to change yet

Do not touch the auth backend or session cookie strategy.

7.3 Dashboard / Trips home

User intent

See my trips, jump back in, plan the next one.

Visual goal

An editorial reading room: a short concierge greeting, two or three trip "covers" with destination atmosphere, a quiet AI Concierge launchpad.

Information hierarchy

- 1. Concierge greeting line ("Good evening, Prashanth. Three trips on your shelf.") — serif Display M.
- 2. Trip covers (large, full-bleed photo + cover details).
- 3. AI Concierge launch pill, sticky bottom on mobile.
- 4. Travel hints strip ("Your Chicago trip is in 18 days — restaurants are filling Saturdays.")

Key components

- Trip cover card (see §8.x).
- Greeting block (h1) with destination memory pill.
- Concierge launch pill (sticky on mobile, sidebar entry on desktop).
- Hint strip with at most 3 hints, all ranked by trip recency.

Interactions

- Tap a trip cover to open trip detail with the cover image cross-fade (see §5.2).
- Tap the concierge pill to open the drawer with the trip in context.

Empty / loading / error

- Empty: a single typographic state — "No trips yet. Where to next?" — with a single CTA. No illustration.
- Loading: the greeting paints immediately with the user's name; trip covers fade in with stagger.
- Error on trip fetch: a calm strip at top: "Could not load your trips. Retrying." with a manual retry link.

Risks

- Trip cover photos require sourcing — if we do not have one, we typeset the destination name in serif. Never stock photo.

Implementation notes

- Reuse existing /trips data; no API change.

GUARDRAIL · What not to change yet

Do not introduce a new trip-creation flow on this page; the "plan a trip" CTA opens the existing flow.

7.4 Trip detail page

User intent

Plan, refine, and live in this trip.

Visual goal

A reading-room layout with three persistent regions: cover + memory header, body (tabs: Itinerary / Saved Ideas / Explore), AI Concierge drawer that is always one keystroke away.

Information hierarchy

- 1. Trip cover: full-bleed image + destination + dates + party size.
- 2. Tabs: Itinerary (default), Saved Ideas, Explore, Notes.
- 3. Right-rail summary on desktop: hotel, flights, total saved spots.
- 4. AI Concierge drawer (collapsed by default).

Key components

- Trip cover header, Tab bar, Itinerary timeline, Saved Ideas grid, Explore grid, Concierge drawer toggle.

Interactions

- Tabs switch with content cross-fade.
- Drag-to-day from Saved Ideas to Itinerary.
- Concierge drawer remembers trip context (memory pill).

Empty / loading / error

- Empty itinerary: typographic, with three concierge prompts ("Restaurants we should book first?", "Best half-day in Lincoln Park?", "Anything unmissable for an anniversary?").
- Loading: the cover paints immediately; tabs and itinerary fade in.
- Error: per-tab error chip, not a page-level red banner.

Risks

- Drag interactions on mobile must be long-press to avoid scroll conflicts.

Implementation notes

- Existing trip detail components are mostly intact; this is a visual + IA pass, not a data refactor.

GUARDRAIL · What not to change yet

Do not change the trip data model or any persistence contract.

7.5 AI Concierge drawer / page

User intent

Ask the concierge anything; get verified, addable answers.

Visual goal

Card-first, see §6 for the full chapter.

Information hierarchy

1. Memory pill. 2. Intent line. 3. Result grid. 4. Refinement chips. 5. Composer.

Key components

- Memory pill, intent line, result grid (§8), thinking breadcrumb, refinement chips, composer.

Interactions

See §6 in full; this entry is a stub for completeness.

Empty / loading / error

See §6.7–6.8.

GUARDRAIL · What not to change yet

Do not modify intent detection, query parsing, ranking, or evidence pipelines. Frontend only.

7.6 Explore tabs (Areas, Restaurants, Attractions, Nightlife, Hotels)

User intent

Browse curated, verified options when I do not have a specific question.

Visual goal

A magazine-section feel, with tabs as discreet metadata rather than dominant chrome. Each tab uses the same Card primitive but with a distinct top motif (a sandstone hairline + tab-name overline) to provide rhythm without changing the card itself.

Information hierarchy

- 1. Section title ("Restaurants in Chicago") + filter chip row.
- 2. Map ↔ list split (desktop) or list-default (mobile).
- 3. Pagination/lazy load — not infinite scroll without anchor.

Key components

- Tab bar (Areas / Restaurants / Attractions / Nightlife / Hotels).
- Filter chip row (cuisine, price band, neighborhood, vibe).

- Map (themed per §4.9).
- Card grid (3 desktop / 2 tablet / 1 mobile).

Interactions

- Filter chips persist in URL state (Expedia pattern, §2).
- Map ↔ card hover sync (§5.7).
- "Search this area" chip on pan; never auto-requery.

Empty / loading / error

- Empty: typed, with three refinement chips.
- Loading: 9 skeleton cards in dark mode.
- Error: per-section coral hairline strip.

Risks

- Hotels tab depends on hotel data we may not have for every destination; design must handle "no inventory yet" gracefully.

Implementation notes

- Areas tab uses neighborhood polygons styled per §4.9; if absent, fall back to typeset neighborhood list.

GUARDRAIL · What not to change yet

Do not introduce hotel booking flow yet; if Hotels tab opens, it is browse-only with a hand-off link.

7.7 Saved Ideas / Trip Ideas

User intent

A scrapbook of places I might want, with the lightest possible commitment.

Visual goal

Light-mode, paper surface. The only place we use light mode *inside* the app shell, because saved ideas are the trip's artefacts. The shift in tone is intentional and signals "this is yours, keep it."

Information hierarchy

1. Date-grouped sections ("Saved this week", "From your Paris trip", etc.).
2. Card grid in paper-mode treatment.
3. Drag handles always visible on desktop.

Key components

- Saved Idea card (paper variant of §8 card).
- Section headers in serif Display S.

- "Move to trip" inline action per card.

Interactions

- Drag-to-trip on desktop, long-press-to-move on mobile.
- Bulk select with keyboard (J/K to navigate, X to add to bin).

Empty / loading / error

- Empty: "Your scrapbook is empty. Tap ♥ on anything to keep it."
- Loading: skeletons in paper tone.
- Error: hairline coral strip with a manual retry.

Risks

- The dark-to-light tonal switch must be done at the page-template level, not per-component, to avoid collision.

Implementation notes

- One CSS class on the route container drives the mode switch; tokens do the rest.

GUARDRAIL · What not to change yet

Do not split saved ideas storage from trip-scoped saves; the data model stays.

7.8 Itinerary timeline / day planning

User intent

See and shape my day-by-day plan.

Visual goal

Editorial timeline. Days are not boxes; they are typeset sections separated by a hairline. Items are cards. Travel-time hints are caption-grade typography between cards.

Information hierarchy

- 1. Trip dates header (always visible).
- 2. Day section: serif Display S date + meta ("Sat, May 9 · light rain forecast").
- 3. Itinerary cards in chronological order.
- 4. Travel-time hints between cards.

Key components

- Itinerary card (§8), Day section header, Travel-time hint, AI Concierge "Suggest for this day" inline trigger.

Interactions

- Drag to reorder; drop zones span the full row.

- Tap a card to open its detail drawer.
- "Suggest for this day" opens the Concierge with the day in context.

Empty / loading / error

- Empty day: a single typographic line ("Nothing planned for Saturday yet.") and a "Suggest for this day" chip.
- Loading: the day headers paint immediately; cards stagger in.
- Error: day-level coral hairline strip.

Risks

- Drag-and-drop reorder must work with keyboard for accessibility.

Implementation notes

- Reuse existing itinerary state; no API change.

GUARDRAIL · What not to change yet

Do not introduce time-of-day blocks (morning/afternoon/evening) on top of explicit times unless the user explicitly opts in.

7.9 Travel hints, flights, hotels, activities

User intent

See research-stage information without committing.

Visual goal

Caption-grade type for hints; cards reused for hotels/flights/activities, with the same trust strip discipline.

Information hierarchy

- Hints are always inline with the relevant itinerary item.
- Hotels/flights/activities are full sections in Trip Detail with the same card grid as Explore.

Key components

- Travel-time hint (caption + chip "view route").
- Hotel card variant of §8.
- Flight card variant of §8 (price, stops, duration).

Interactions

- "View route" opens the map with the route plotted.
- "Open in booking provider" hands off; we do not book in-app yet.

Empty / loading / error

- Empty: section omitted entirely; we do not show "No hotels yet" placeholders.

Risks

- Provider hand-off must use a clean target=_blank with rel="noopener".

Implementation notes

- No new providers in design slices; just visual reuse.

GUARDRAIL · What not to change yet

Do not introduce in-app booking until a separate, explicit roadmap phase.

7.10 Cards and modals

User intent

Drill into a place, save it, add it, share it.

Visual goal

A half-bleed detail panel (right rail on desktop, full sheet on mobile) rather than a full-screen modal. The panel breathes; the page stays visible behind it dimmed to 60%.

Information hierarchy

1. Hero photo (or typeset monogram).
2. Place name (Display M serif), trust strip, neighborhood + walking distance.
3. "Why this fits" pull-quote with sources.
4. Metadata grid: hours, price band, phone, website.
5. Map preview with pin.
6. Actions: Add to trip, Save, Share.

Key components

- Card detail panel, Trust strip, Pull-quote, Metadata grid, Map preview.

Interactions

- Esc closes; backdrop click closes; deep-link with URL param.

Empty / loading / error

- Loading: hero placeholder, content paints first.
- Error: panel header + "Could not load this place" + retry.

Risks

- Half-bleed panels can collide with mobile bottom sheets; on mobile, panel becomes a 90vh sheet identical to AI Concierge.

Implementation notes

- Reuse Concierge sheet shell; one panel primitive for both surfaces.

GUARDRAIL · What not to change yet

Do not introduce in-card booking actions yet.

7.11 Profile / settings

User intent

Manage account, preferences, and travel style.

Visual goal

Light mode, two-column on desktop, single-column on mobile. Sections are paper cards. Restraint over breadth.

Information hierarchy

- 1. Identity: name, email, photo (optional), default home city.
- 2. Travel style: party size default, pace (slow/balanced/fast), cuisine likes/dislikes.
- 3. Account: password, sign out.
- 4. Privacy: data export, delete account.

Key components

- Section card, Inline-edit field, Travel-style chip group.

Interactions

- Inline edit; save on blur or enter; never a "Save profile" big button at the bottom.

Empty / loading / error

- Loading: section cards skeleton; identity paints first.

Risks

- Travel-style preferences may eventually feed the AI Concierge; the design must accommodate that without making the UI noisy now.

Implementation notes

- No new fields shipped beyond what the backend supports today.

GUARDRAIL · What not to change yet

Do not wire travel-style preferences into the Concierge in this design slice.

7.12 Mobile nav / sidebar

User intent

Move between trips, dashboard, and the Concierge.

Visual goal

A bottom rail of three icons + concierge pill on mobile. A slim 72px sidebar on desktop with iconography only, expanding to 240px on hover/click.

Information hierarchy

- 1. Dashboard. 2. Trips. 3. Concierge (pill). 4. Profile.

Key components

- Bottom rail, Sidebar, Concierge pill (sticky).

Interactions

- Tap rail items to navigate; Concierge pill always opens the sheet/drawer in context.

Empty / loading / error

- No empty state.

Risks

- Bottom rail collides with iOS home-bar; pad 12px.

Implementation notes

- The sidebar shipped in PR #168 already; this is a refinement, not a rebuild.

GUARDRAIL · What not to change yet

Do not change navigation routes.

SECTION 08

Card design system

Cards are the atomic unit of the product. One primitive, one shape language, configurable slots.

8.1 The Card primitive

There is one Card primitive. Every other card type composes it with different slots. The primitive enforces tokens; variants only choose which slots to render.

Slot	Required	Notes
Identity (name + place type)	always	Display S serif on detail; Body L on grid.
Trust strip	always	Verified mark + source count + confidence chip.
Media (image or monogram)	optional	Aspect 4:5 grid, 3:2 detail; never stock.
Why this fits	optional	Pull-quote with at least one source tag.
Metadata strip	optional	Neighborhood · price band · walking time.
Actions	always	"Add to trip", "Save", "Open" affordances.
Caveat	conditional	Weak-evidence sentence in caution amber.

8.2 Card variants

Variant	Used in	Slot config	Key cue
Verified place card	AI Concierge, Explore	identity, trust, media, why, meta, actions	Sandstone "Verified by Google" mark.
AI Concierge result	Concierge result canvas	identity, trust, why, meta, actions; media is hero	Why-this-fits is hero typography.
Explore card	Explore tabs	identity, trust, media, meta, actions	Top hairline + tab overline (cuisine/area).
Saved idea card	Saved Ideas (paper)	identity, media, meta, action: move-to-trip	Paper surface; warm tones.
Itinerary item card	Itinerary timeline	identity, time, meta, action: edit/move	Time stamp typography is the lead.
Area card	Areas tab	identity, neighborhood polygon thumbnail, neighborhood map preview	Neighborhood map preview replaces hero.
Hotel card	Hotels tab	identity, trust, media, price band, actions	Price band is part of meta strip.
Flight card	Flights research	identity, route, time, price, actions	Typographic route ribbon (DEN → ORD)
Activity card	Activities research	identity, trust, media, duration, price, actions	Duration chip in caution amber if >1 day

8.3 Trust badges

- **"Verified by Google".** — A small icon + caps overline label, sage on dark, sage on paper. Only on cards where backend has a stable Google Place ID and OPERATIONAL status.
- **"3 sources".** — A typographic chip rendered as the count + the word "sources". Tapping expands the source list inline.
- **"Confidence: high / medium / low".** — Plain-English chip mapped to evidence count. Never a numeric percentile.

- **"Worth the splurge" / "Luxury for less".** — See §3.7. Brass background, ink text. Used selectively, not on every card.

8.4 Reason / why sections

On every card variant, the "why this fits" string is the one we receive from the backend (e.g., `_build_dynamic_why()`). The UI never paraphrases it. It is rendered as a serif italic pull-quote (Quote style, §4.4) with an inline source tag at the end. If the reason is generic, the source tag reads (*verified place, no editorial*).

8.5 Source / evidence indicators

- **Source list is one tap away.** — Tapping the source count reveals a typeset list of sources (publication name, date, link). The list is not collapsed in HTML — it lives inline so screen readers see it.
- **Source chips, not stars.** — We never aggregate sources into a star rating. If aggregation is meaningful (e.g., 3 of 4 mention "tapas"), we render: *"3 of 4 sources call out tapas/small plates."*
- **Date discipline.** — Sources show a date when available. If a source is > 24 months old, it gets a faint "older source" caption.

8.6 Add / save actions

- **Primary: "Add to trip".** — Sandstone gold filled button on dark; brass on paper. 36px height on grid, 44px in detail panel. Always paired with a label.
- **Secondary: heart-toggle save.** — Outline → fill on save. No "Saved!" toast; sidebar dot pulses (see §5.5).
- **Tertiary: "Open" / "Share".** — Hairline outline buttons. "Open" handsoff to map or external; "Share" copies a shareable link in light-mode paper aesthetic.

8.7 Map-linked behavior

- **Hover to highlight.** — Hovering a card highlights its pin and dims unrelated pins (§5.7).
- **Tap pin to open card detail.** — Same affordance as tapping the card itself; deep-linked.
- **Travel-time chip.** — When in trip context, every card surfaces "12 min walk from your hotel" or "22 min from the Field Museum (your last stop)" if backend provides it. When not in trip context, this chip is omitted — never invented.

8.8 Compact vs expanded variants

Each card has two density modes: **compact** for grid contexts, and **expanded** for the detail drawer. The compact mode hides the why-pull-quote behind a "show why" link; the expanded mode renders it as the hero block. The trust strip is identical in both modes — never compressed into a single icon.

8.9 What the card system forbids

REJECT · Card anti-patterns

No price-strikethrough. We are not an aggregator.

No urgency badges. No "limited", "popular", "going fast".

No "edited by AI" labels. The card is the AI's output; if it is on screen, it is verified. Labels imply doubt.

No multi-color category badges. Categories live in the meta strip as caption text, not as color-coded pills.

No "Recommended for you" pseudo-personalization. We use the user's real preferences when we have them, or we say nothing.

SECTION 09

Evidence and trust UX

How evidence appears beautifully, honestly, and without fake precision.

Evidence is the brand. Other AI travel planners hide their reasoning behind fluent prose; Travel Concierge surfaces it. The design challenge is to make evidence look like editorial pull-quotes, not legal disclaimers.

9.1 The evidence ladder

A four-rung ladder maps backend confidence to a UI treatment. The ladder is the only mapping; no surface invents its own.

Rung	Backend signal	UI treatment	Voice example
Verified place + 3+ sources	Google Place ID OPERATIONAL	"Verified by Google" badge, sage source chip, and multiple dates	"Angelo's is a past favorite match than the..."
Verified place + 2 sources	Google Place ID OPERATIONAL	"Verified by Google" badge, sage source chip	"For the trip brief with a dinner-date feel in..."
Verified place, weak constraints	Google Place ID OPERATIONAL	"Verified by Google" badge, sage source chip, (Best trip constraints) list with polished de	
Place not verified	No stable Google Place ID, or NO OPERATIONAL	"Not Verified" badge; treated as editorial reference only	"Mentioned in two recent guides; we could n..."

9.2 "Verified by Google" mark

- **Where it sits.** — Top-left of every verified card, above the place name. Caption-grade typography.
- **What it looks like.** — A 12px sage filled circle + 8px sage caps overline "VERIFIED BY GOOGLE". Not an icon-only badge; the words matter.
- **What it never means.** — It does not mean "we recommend this place". It does not mean "Google recommends this place". It means: a stable Google Place ID exists and the business is OPERATIONAL.

9.3 Source counts

- **Display.** — "3 sources" rendered as a typographic chip with a hairline outline. Tap to expand the inline source list.
- **Source list.** — Each row: publication name + date + 1-line snippet that supports the claim. Snippets are extractive, never paraphrased.
- **Threshold.** — We display the count starting at 1. We do not hide low source counts; transparency over flattering numbers.

9.4 Evidence quality (confidence chip)

A second small chip: *Confidence: high / medium / low*. It is plain English. We never show a number like "92% confidence" because we cannot defend the precision.

Chip	Backend rule (illustrative)	Color
------	-----------------------------	-------

Confidence: high	≥ 3 sources + place verified + constraint corroborated	sage
Confidence: medium	2 sources + place verified	sandstone gold (no fill)
Confidence: low	1 source + place verified, OR ≥ 1 unverified constraint	caution amber

9.5 "Why it fits" pull-quotes

The why-quote is the editorial heart of the card. It is rendered in serif italic, 16px on grid / 18px in detail. It uses curly quotes. It cites a source tag. It never paraphrases the backend reason; it shows it verbatim with hairline-styled quotation marks.

“A stronger tapas/small-plates match than a generic cocktail bar in West Loop.”
— verified place, dynamic match

9.6 "Best for" tags

A small caps overline above the place name: *BEST FOR · DINNER DATES*, *BEST FOR · GROUPS OF 4-6*. These tags come only from corroborated evidence; we never invent them. If we have no evidence to support a "best for" tag, we omit the slot entirely.

9.7 Caveats

The caveat slot is a *typeset sentence*, not an icon. It uses caution amber for the leading word and ink/pearl for the rest of the line. Examples:

- **"Verify when booking"** — — we could not confirm the waterfront view.
- **"Older sources"** — — two reviews are from 2022; menu may have shifted.
- **"Limited evidence"** — — one independent source; we trust the verification, less so the vibe.

9.8 Weak evidence

OPPORTUNITY · How to communicate weakness as a feature

Render the card. Do not hide it. The user wants to know what is close to the brief.

Promote the caveat. The amber sentence sits directly under the trust strip, before the why-quote. It earns its place.

Reduce the why-quote prominence. Render at compact size; do not promote it to detail-panel hero.

Offer a refinement chip. "Find places with stronger waterfront evidence" appears below the card, generated from the negative constraint.

9.9 Missing details

When a metadata field is unknown (price band, walking distance, hours), we omit the field entirely. We never show "—", "N/A", "TBD", or a placeholder. The absence of the field is the design statement.

9.10 Unsupported claims

REJECT · What we forbid the UI from rendering

Awards. Michelin stars, James Beard, Bib Gourmand — never rendered unless the backend has a structured, dated source.

"Stunning views" / "Romantic ambiance". Subjective adjectives derived from anything other than a corroborated, verbatim source quote.

Distance / walking time. Rendered only when the backend computed it from real coordinates — never estimated client-side.

Neighborhood claim. Rendered only when Google Places has a neighborhood label or the source explicitly names it.

Price band. Rendered only from Google price level. We never infer price from a venue name or photo.

Opening hours / "open now". Rendered only from a fresh Google response. Stale hours are dropped, not displayed.

9.11 Confidence without fake precision

No percentages. No 5-star aggregate scores. No relative ranking like "#1 of 487". Three plain-English confidence rungs, source counts, and extractive snippets — nothing else. The product's premium feel comes from typographic restraint, not from manufactured precision.

SECTION 10

Addictive personalization ideas

Tasteful, evidence-grounded details that make the user return — without urgency, gimmicks, or childish gamification.

These ideas are all opt-in by behaviour — each only appears when the user has supplied enough data for it to be honest. None of them require new fabrication, new claims, or new ranking magic. Each is implementable as a frontend pass over data the backend already exposes.

Evolving trip covers

Each trip cover is a real photo from a verified place inside the trip, rotating to a new one as the trip evolves. If the user adds a hotel in River North, the cover swaps to the hotel hero. If they add a sunrise activity, the cover swaps to that. Never stock; only verified photos.

Destination atmosphere

The trip detail page shifts a single tonal accent (a sandstone vs. brass mix) to match the destination's curated atmosphere — "northern, cooler" vs. "tropical, warmer". Subtle: a 4%-luminance shift, not a full re-skin.

Personal travel style memory

After three saves, the AI Concierge surfaces a one-line summary: "You tend to save quieter, smaller, dinner-date places. Want me to default to that vibe?" Yes/no. The choice persists. We never auto-default without asking.

Luxury-for-less picks

On any Explore tab, a single "Luxury for less" rail at the top, capped at 3 cards. Inclusion requires backend evidence that the place delivers a high-end experience at a notably lower price band than peers. Never a "deal of the day".

Worth-the-splurge moments

On any verified card, when the place is meaningfully higher-priced AND has corroborating evidence of a notable experience, a brass "Worth the splurge" chip appears next to the price band. Selectively. Never on more than ~10% of cards in a result set.

Rainy-day swaps

When a trip is < 7 days away and weather data shows rain on a planned day, a Concierge prompt appears: "It is forecast to rain Saturday. Want me to suggest indoor alternatives for your 2 PM block?" Honest, optional, never automatic.

"Tonight's edit"

On the dashboard, a single inline strip: "Three places open tonight that fit your saved style." Only appears when the user is in the destination city (geo permission granted) and the data is fresh.

Neighborhood mood maps

On the Areas tab, each neighborhood polygon tints with a tonal accent corresponding to its dominant character (food, nightlife, museums). The tint is computed from the count and category of verified places in the polygon, not editorial.

Saved-card collections

Manual collections inside saved ideas: "Anniversary picks", "Worth a flight". Drag-and-drop. No suggested collections; the user names them.

Trip completion rituals

After the last day of a trip, the dashboard shows a single typographic card: "Chicago, May 5–9. 14 places saved, 9 visited, 3 still in your shelf." A small "save trip to scrapbook" link archives it. No badge, no celebration animation.

Subtle unlocks / progress

Light progress hints: "You have planned 3 of your 4 days." A hairline progress glyph in the trip cover, never a percentage ring.

Conversational next-best-actions

After every Concierge result, a single line: "Want me to also find a wine bar nearby for after?" derived from the ask. Single-tap continues the chain.

10.1 What we never do

REJECT · Personalization anti-patterns

Streaks, XP, levels, badges. Childish; antithetical to the brand.

"Limited time" / "ending soon" prompts. Manufactured urgency.

Animated success bursts. See §5.9.

"You unlocked X" modals. The user did not unlock anything; we just showed them more of what they already chose.

Push notifications about unrelated trips. Email or in-app, never noisy cross-trip pings.

Auto-defaults without asking. Memory is consensual.

SECTION 11

Implementation roadmap

A phased redesign that does not break existing functionality. Each phase is one pull request, one merge gate, one rollback path.

Each phase is independently shippable, independently rollbackable, and bounded by the UI budget gate in docs/ai/PROMPT_LIBRARY.md. No phase is allowed to start until the previous one has merged and survived a 48-hour stability window.

Phase 0 — Design tokens & foundation

Scope. Add CSS variable tokens for colors, type, spacing, elevation, and motion; introduce the Card primitive shell (no variants wired yet); ensure Tailwind theme reads from CSS variables. Touch ≤ 6 files, all in *frontend/*.

Field	Value
Recommended model	Sonnet (precise mechanical work; Codex if PR < 200 LOC).
Risk	Low — tokens replace existing utility classes incrementally.
Frontend tests / screenshots	Visual diff of Dashboard, Trip Detail, AI Concierge drawer at desktop + mobile; reduced-motion test; WC.
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR; tokens are additive on top of existing styles.

HANDOFF.md / progress_log.md update. HANDOFF.md gets a "Design Foundation Phase 0" entry naming the new token files; progress_log.md notes the merge and the 48-hour stability window.

Phase 1 — Shell & navigation polish

Scope. Sidebar density, mobile bottom rail polish, page header consistency, page transitions per §5.2. Reuse existing routes; no IA changes. Touch ≤ 6 files.

Field	Value
Recommended model	Sonnet.
Risk	Low.
Frontend tests / screenshots	Manual nav test on desktop + mobile; reduced-motion test; tab focus order.
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR; no schema or API changes to undo.

HANDOFF.md / progress_log.md update. Note: shell now consumes Phase 0 tokens; legacy classes removed from header and sidebar.

Phase 2 — AI Concierge flagship redesign

Scope. Apply §6 in full: card-first hierarchy, conversation rail, memory pill, composer, thinking breadcrumb. Frontend only; no change to *fast_dynamic_place_search* or routes/ai.

Field	Value
Recommended model	Sonnet (high care).
Risk	Medium — the flagship surface; visual-regression risk is meaningful. Bound to ≤ 8 files.
Frontend tests / screenshots	Manual: tapas-bar query, sushi-with-waterfront query, compare-top-3, more-options pool hit, weak-evidence
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR; old AI Concierge component is preserved as <code><i>AIConcierge.legacy.tsx</i></code> for one release.

HANDOFF.md / progress_log.md update. HANDOFF.md adds "AI Concierge Visual Pass v1" with the screenshot matrix link; progress_log.md notes the merge.

Phase 3 — Card system redesign

Scope. Implement §8 in full: one Card primitive + 9 variants. Replace existing card components route by route, starting with verified place card and AI Concierge result card. Touch ≤ 6 files per slice; this phase may need 2 PRs.

Field	Value
Recommended model	Sonnet, with Codex visual merge gate.
Risk	Medium — cards are the atomic unit; regressions ripple. Mitigation: variants are added behind a fe
Frontend tests / screenshots	Visual diff of every surface that uses a card; trust strip rendering; weak-evidence rendering; add-to-trip a
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Per-surface fallback flag flips, then revert PR.

HANDOFF.md / progress_log.md update. HANDOFF.md adds "Card Shell v1" entry with variant list.

Phase 4 — Explore redesign

Scope. Apply §7.6 to all five Explore tabs. Adopt the Phase 3 card variants. Add filter chip URL persistence. Map theme per §4.9. Touch ≤ 6 files.

Field	Value
Recommended model	Sonnet.
Risk	Medium — map styling can break with provider quirks.

Frontend tests / screenshots	Filter persistence across reload; map theme on light/dark; cluster behavior; chip-derived empty state.
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR; map theme falls back to Google default.

HANDOFF.md / progress_log.md update. HANDOFF.md adds "Explore Visual Pass v1".

Phase 5 — Itinerary & timeline redesign

Scope. Apply §7.8: editorial day sections, hairline travel-time hints, drop-zone affordances, day collapse. Touch ≤ 5 files.

Field	Value
Recommended model	Sonnet.
Risk	Medium — drag-and-drop a11y must be preserved.
Frontend tests / screenshots	Keyboard reorder; touch long-press reorder; reduced-motion; collapsed-day expand; travel-time hint exp
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR.

HANDOFF.md / progress_log.md update. HANDOFF.md adds "Itinerary Visual Pass v1".

Phase 6 — Landing & auth redesign

Scope. Apply §7.1 and §7.2. Cinematic hero (still photo + 1% grain), glass composer pill (only sanctioned glass surface), still auth photo. Two PRs may be cleaner: hero first, auth second.

Field	Value
Recommended model	Sonnet, with Codex visual merge gate.
Risk	Medium — first impression; mobile FCP must be guarded.
Frontend tests / screenshots	Lighthouse FCP ≤ 2.0 s on 3G Fast emulation; reduced-motion; image LQIP fallback test.
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR; existing landing component preserved as <code><i>landing.legacy.tsx</i></code> .

HANDOFF.md / progress_log.md update. HANDOFF.md adds "Landing & Auth Visual Pass v1".

Phase 7 — Mobile refinement

Scope. Sweep all surfaces for mobile thumb reach, sheet behaviors, AI Concierge sheet, Itinerary touch targets. Touch ≤ 6 files.

Field	Value
Recommended model	Sonnet.
Risk	Low.
Frontend tests / screenshots	Manual phone tests on iOS Safari + Android Chrome.
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR.

HANDOFF.md / progress_log.md update. HANDOFF.md adds "Mobile Refinement v1".

Phase 8 — Motion polish

Scope. Apply §5 in full: page transitions, card entrance, thinking breadcrumb, save/add micro-moments, reduced-motion. Bounded to motion code only; no layout changes.

Field	Value
Recommended model	Sonnet, with Codex visual merge gate.
Risk	Medium — motion regressions are subtle.
Frontend tests / screenshots	Reduced-motion test; FCP regression check; battery sanity on phone for 90 s of normal use.
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR; motion tokens fall back to instant.

HANDOFF.md / progress_log.md update. HANDOFF.md adds "Motion Polish v1".

Phase 9 — Final cohesion pass

Scope. Sweep for stragglers, light-era classes, inconsistencies discovered post-Phase 8. Capped to 6 files of removals/unifications. No new components.

Field	Value
Recommended model	Codex.
Risk	Low.

Frontend tests / screenshots	Visual diff per surface against the Phase 0 baseline.
Backend changes allowed?	No.
SQL allowed?	No.
Rollback strategy	Revert PR.

HANDOFF.md / progress_log.md update. HANDOFF.md adds "Design Bible Cohesion Pass — complete".

11.x Roadmap shape

PRINCIPLE · Why these phases, in this order

Foundation first. Phase 0 is the only phase that adds tokens. Every other phase consumes them.

Flagship before breadth. Phase 2 (AI Concierge) precedes Explore and Itinerary because the Concierge sets the visual bar.

Cards before surfaces that use them. Phase 3 precedes 4, 5 because every later phase reuses the card primitive.

Landing last among visual. First impression edits last so the bar is set before we polish the welcome mat.

Motion last. Motion can mask layout problems; we want layout right first.

No backend, no SQL, anywhere. If a phase needs them, it is split before it starts — never bundled into a design PR.

SECTION 12

Guardrails for future implementation prompts

The rules every Sonnet/Codex design prompt must follow. These are the diff between a great visual pass and a regression.

12.1 Hard rules

GUARDRAIL · Non-negotiable guardrails for every design PR

- 1. Preserve payload contracts.** A design PR may not change any field on the AI Concierge response, the place card payload, the trip payload, or the itinerary payload. Adding a new optional UI-only field is allowed only if it is computed from existing fields entirely on the client.
- 2. Preserve addable card behavior.** Verification gates, identity keys, pool/dedup, and the "no fake addable cards" invariant are sacred. Cards render only when the backend marked them OPERATIONAL with a stable Place ID.
- 3. Preserve tests.** Every existing backend and frontend test must pass. A failing test is a stop. We never delete tests to make a UI PR pass.
- 4. No backend rewrites.** A design PR may not modify *backend/app/services/**, *routes/**, *concierge/**, or any ranking, parsing, or evidence pipeline. If a UI claim requires backend data, the data is fetched in a separately scoped backend PR *first*.
- 5. No SQL.** A design PR may not include Supabase migrations. The PR summary must explicitly state "Supabase SQL: No". If the design appears to need SQL, the PR is rejected and the work is split.
- 6. No intelligence behavior changes.** A design PR may not edit prompts, intent detection, query parsing, ranking, scoring, or category gates. If a visual treatment depends on these, the dependency goes through a separate backend PR.
- 7. No hidden errors.** If a fetch fails, the UI shows a calm error with a retry. We never swallow errors to make a UI look stable.
- 8. No fake trust signals.** A "Verified by Google" mark may render only when the backend payload says the place is OPERATIONAL with a stable Place ID. Identical for source counts and confidence chips.
- 9. No large all-in-one redesign PR.** Every PR maps to exactly one roadmap phase (§11). A PR that touches more than 6 frontend files needs Code Committee approval before merge.
- 10. Reduced-motion test required.** Every UI PR must include a *prefers-reduced-motion: reduce* verification screenshot or note.

12.2 PR prompt template (mandatory)

Future Sonnet/Codex design prompts must conform to this template. The template is short by design.

Title: "Design Phase X — [scope]"

Sections cited: §11 Phase X, §[other relevant chapters]

Files allowed: <explicit list, ≤ 6 unless approved>

Files forbidden: *backend/***, *supabase/***, *routes/***, *services/***

Frontend tests: <list of manual smoke checks>

Reduced-motion test: screenshot or note

Supabase SQL: No

Rollback: Revert PR; no schema or API changes to undo.

HANDOFF.md update: [yes/no + 1 line]

12.3 Soft rules (strong defaults)

- **Use tokens, never hex.** — If a token is missing, add the token in a separate Phase 0 follow-up PR first.
- **Use the Card primitive.** — No new bespoke card components without a Code Committee review.
- **Use the existing motion easings.** — No new easing curves; pick from the four defined in §5.
- **Run the spouse-friendly smoke test.** — After every UI PR, manually run: open dashboard → tap a trip → ask the concierge for tapas → save a card → see it in the trip. Mobile + desktop. Note any regressions in the PR.
- **Cap one design PR per branch per day.** — Avoids stacked-PR review fatigue and CI flakiness.

12.4 Stop conditions

GUARDRAIL · When to stop the PR and reclassify

The PR exceeds 6 files. Stop, split.

The PR needs a backend change. Stop, split into a backend PR first.

A test fails. Stop. Diagnose. Do not delete the test.

The Concierge returns 0 cards on the smoke test after the PR. Stop. Roll back. Diagnose offline.

FCP regresses on landing or Trip Detail. Stop. Roll back.

Reduced-motion mode is not tested. Stop. Add the test before merging.



SECTION 13

First recommended design implementation slice

A small, safe, foundation-level slice that establishes the visual system without touching intelligence or content.

13.1 Recommendation

The first PR is **Phase 0 — Design tokens + Card primitive shell**, and nothing else. It is deliberately small enough to merge in a single Sonnet pass, deliberately important enough that every subsequent phase inherits its structure.

13.2 Why this slice

- **Foundational.** — Every later phase consumes the tokens and the Card primitive.
- **Reversible.** — Tokens are additive; reverting the PR removes them without touching consumers.
- **Visual but invisible.** — The user sees no immediate change; downstream phases unlock the visible work.
- **Avoids product risk.** — No surface adopts the new card variants in this PR; we cannot regress rendering of any card we do not touch.

13.3 Scope (precise)

- **Add tokens to *frontend/src/app/globals.css*:** — colors (12 dark + 9 paper), type scale (11 roles), spacing (10 steps), elevation (5 levels), motion (4 easings, 4 durations), per §4.
- **Wire Tailwind theme to read CSS variables.** — No raw color values inside Tailwind config — only `var(--token)`.
- **Add a Card primitive at *frontend/src/components/ui/Card.tsx*:** — composable slot API (identity, trust, media, why, meta, actions, caveat) with no variants wired yet. Pure structural component.
- **Add a TrustStrip primitive at *frontend/src/components/ui/TrustStrip.tsx*:** — renders the verified mark, source count, confidence chip from a simple props shape that mirrors existing payloads.
- **Update *UI_BASELINE.md*** — Note that token foundation v2 has shipped and link to this bible.

13.4 Out of scope (explicitly)

GUARDRAIL · What this slice does NOT do

No surface adopts the new Card primitive in this PR.

No backend or API change. No SQL.

No motion implementations beyond declaring the duration and easing tokens.

No new icon set. No new font hosting. No new animation library.

No removal of existing components yet — everything coexists.

13.5 Files allowed (≤ 6)

#	File	Change
1	<code><i>frontend/src/app/globals.css</i></code>	Add CSS variables for color, type, spacing, elevation, motion.
2	<code><i>frontend/tailwind.config.ts</i></code>	Theme keys read from <code>var(--token)</code> .
3	<code><i>frontend/src/components/ui/Card.tsx</i></code>	New composable Card primitive.
4	<code><i>frontend/src/components/ui/TrustStrip.tsx</i></code>	New TrustStrip primitive.
5	<code><i>docs/ai/UI_BASELINE.md</i></code>	Note token foundation v2 + link to bible.
6	<code><i>docs/ai/HANDOFF.md</i></code>	Add "Design Foundation Phase 0" entry.

13.6 Acceptance checks

- **Build passes; type checks pass.** — Existing UI is visually unchanged.
- **Tokens visible in DevTools.** — Hover any element — computed colors show `var(--surface-1)` not a raw hex.
- **Card primitive renders an empty shell.** — A Storybook entry or a `/debug` route demonstrates the Card primitive in dark and paper modes with each slot toggled.
- **Reduced-motion respected.** — Card primitive has no entrance animation by default.
- **Spouse-friendly smoke test.** — Dashboard → trip → AI Concierge → tapas search → add to trip works identically before and after the PR.

13.7 Rollback

Single PR revert. Tokens removed; Tailwind theme returns to its prior shape; two new components are removed. No data, no schema, no API to unwind.

13.8 What is NOT yet a prompt

PRINCIPLE · No implementation prompt yet, by design

Per the brief, this bible recommends the first slice but does not author the implementation prompt. Pressure-test this bible (especially §6, §8, §9, §11) in ChatGPT or with a second reviewer first. After agreement, write the Phase 0 prompt against §13.3–13.6 verbatim.

13.9 After Phase 0

When Phase 0 is merged and stable for 48 hours, the next slice is Phase 2 (AI Concierge flagship redesign), not Phase 1 (shell polish). The Concierge is the highest-leverage surface; once it consumes tokens via the Card primitive, every other surface follows naturally.

SECTION A

Appendix · Glossary, success criteria, open questions

Definitions and pressure-test prompts for the next review pass.

A.1 Glossary

Term	Meaning in this bible
Verified place card	A card backed by a stable Google Place ID with OPERATIONAL status.
Addable	A card the user can add to a trip; must be a verified place card.
Why-quote	The serif italic pull-quote on every card describing why the place fits.
Trust strip	The horizontal row of trust signals on every card.
Evidence ladder	The four-rung backend-confidence-to-UI mapping in §9.1.
Memory pill	The single editable summary of the active conversational context.
Composer	The AI Concierge prompt input, always visible.
Result canvas	The right-side region of the AI Concierge that shows verified cards.
Conversation rail	The left-side region of the AI Concierge that shows turns + intent.
Tonal duality	The two-mode system: midnight ink for active surfaces, paper for artefacts.
Card primitive	The single composable component that backs every card variant.
Card variant	A configured composition of the Card primitive (see §8.2).
Spouse-friendly smoke test	The end-to-end flow we manually validate after every UI PR.
Roadmap phase	A bounded design PR scope that maps 1:1 to a section in §11.

A.2 Success criteria for the program

- **No design PR regresses backend tests, AI Concierge latency, or addable card count.** —
- **After Phase 3 merges, every surface consumes the Card primitive.** — Audited via grep for raw hex colors and bespoke card components.
- **After Phase 9 merges, the lifecycle UI cost is < 35% of session.** — Measured against the PR #168 baseline of ~51% lifecycle cost.
- **User-reported "this looks like a luxury concierge" sentiment captured.** — Informal but tracked.
- **No fake trust signals shipped.** — Audited via test that asserts Verified mark only renders when the payload flag is present.
- **No invented metadata.** — Audited via test that asserts walking-time, neighborhood, price band fields are absent in the DOM when absent in the payload.

A.3 Open questions for the next review

These are the questions to answer before Phase 0 ships. They are not blockers; they are the inputs to the next decision.

- **Type system.** — Do we self-host an editorial serif (e.g., Editorial New, GT Super, Tiempos) or stay with system serifs for portability and FCP? Recommendation: stay system in Phase 0; add a self-hosted serif in a Phase 6 sub-slice specifically scoped for it.
- **Photography pipeline.** — How do we source verified-place photos at scale without violating provider attribution rules? This determines whether trip covers can use real photographs in Phase 0–3 or remain typeset.
- **Map provider.** — Are we comfortable applying a custom theme to Google Maps, or do we evaluate Mapbox / MapLibre for finer style control? Recommendation: stay on Google with a custom style JSON; defer Mapbox unless a concrete limitation surfaces.
- **AI Concierge memory persistence.** — Currently in-memory. The bible assumes the memory pill summarizes session state only. If we persist across sessions, the pill UX is unchanged but the "forget" link becomes a real durable action; backend out of design scope.
- **Mobile sheet vs. drawer on tablet.** — Tablet viewport is awkward; current bible says sheet. Confirm with a single mobile-tablet review.
- **Light-mode admission.** — Saved Ideas is the only sanctioned light-mode app surface. Confirm with the founder before Phase 5 begins; if rejected, Saved Ideas stays dark and the "artefact" cue is signalled differently.

A.4 Reading order for ChatGPT pressure test

When you bring this bible back to ChatGPT for review, send chapters in this order to maximize useful pushback:

1. **§00 Executive summary** — "challenge any assumption."
2. **§06 AI Concierge UX** — "is the card-first hierarchy correct?"
3. **§09 Evidence and trust UX** — "are the four rungs of the evidence ladder defensible?"
4. **§11 Roadmap** — "is Phase 0 the right first slice?"
5. **§12 Guardrails** — "are these enforceable without slowing the team down?"
6. **§13 First slice** — "what would you cut to make this even smaller?"

A.5 Final check: the mantras

Beauty without logic and logic without beauty are equally bad.

Fast does not mean dumb. Honesty is luxury. Evidence is the brand. The card is the hero. The concierge is composed, not chatty. Every pixel has purpose; every word has a source; every motion has a reason; every PR has a phase.